

# B-VAE Implementation for Disentangled Representations

STD ID: 24081274

## Introduction

This project explores the Beta-Variational Autoencoder ( $\beta$ -VAE), a powerful deep learning model designed to learn disentangled latent representations. It demonstrates how controlled regularization improves interpretability of learned features, enabling better manipulation, visualization, and understanding of complex data structures.

---

## Project Overview

This project focuses on implementing the  $\beta$ -VAE architecture to study and demonstrate disentangled representation learning.  $\beta$ -VAE introduces a hyperparameter  $\beta$  to strengthen the KL-divergence constraint, encouraging the model to separate hidden factors of variation more clearly than standard VAEs. The project includes

- Complete  $\beta$ -VAE implementation in Python
  - Training pipeline with configurable  $\beta$  values
  - Visualization of latent traversals to analyze disentanglement
  - Easy-to-understand code structure suitable for learners and researchers
- 

## GitHub Repository

The full source code, including model architecture, training scripts, and visualization utilities, is available on GitHub:

**GitHub:** [Vishva2003/B-VAE\\_Implementation\\_for\\_Disentangled\\_Representations](https://github.com/Vishva2003/B-VAE_Implementation_for_Disentangled_Representations)

---

## Live Webpage

A demonstration webpage showcasing explanations, concepts, and visual outputs of the  $\beta$ -VAE model is accessible here:

**Webpage:** [beta-vae-tutorial.vercel.app](https://beta-vae-tutorial.vercel.app)

---

## Conclusion

This  $\beta$ -VAE implementation and tutorial webpage together provide an accessible path to understanding disentangled representation learning. The project is suitable for academic submissions, portfolio use, and deeper exploration of generative models.